

CS304 Assignment 4: Implementing TreeSort with a Custom BST
Due: Tuesday March 25th at 11:59 PM (Submit to moodle)

Objective

In this assignment, you will:

- 1. **Implement a Binary Search Tree (BST) from scratch**, including the **Big Five** (Rule of Five).
- 2. **Implement the TreeSort algorithm** using a **in-order traversal** to store sorted elements in a vector.
- 3. **Compare the performance of TreeSort with C++'s std::sort** on large datasets.

Part 1: Understanding TreeSort

How TreeSort Works

TreeSort is a **comparison-based sorting algorithm** that leverages a **Binary Search Tree (BST)**:

- 1. **Insert all elements into a BST** (logically sorting them).
- 2. **Traverse the BST in-order** to extract elements in sorted order.
- 3. **Store the sorted elements in a vector**.

Time Complexity

Operation	Average Case	Worst Case (Unbalanced Tree)
Insertion	$O(\log n)$	$O(n)$
Traversal (Pre-order)	$O(n)$	$O(n)$
Overall Complexity	$O(n \log n)$	$O(n^2)$ (degenerated tree)

♦ **Note:** To improve performance, you could use a **Self-Balancing BST (like AVL or Red-Black Tree)**, but for this assignment, implement a **standard unbalanced BST**.

Part 2: Implementation Requirements

1. Implement a Binary Search Tree (BST) from Scratch

- ✔ You must define a BST class with the following:
 - **Node Structure (TreeNode)**
 - **Insertion Method (insert)**
 - **Pre-order Traversal (preOrderTraversal)**
 - **Destructor, Copy Constructor, Copy Assignment, Move Constructor, Move Assignment (Big Five)**

BST Class Skeleton:

```
template<typename T>
class BST {
private:
    struct TreeNode {
        T data;
        TreeNode* left;
        TreeNode* right;
        TreeNode(const T& value) : data(value), left(nullptr), right(nullptr) {}
    };

    TreeNode* root;

    void insertHelper(TreeNode*& node, const T& value);
    void preOrderTraversalHelper(TreeNode* node, std::vector<T>& sortedVector);
    void destroyTree(TreeNode* node);
    TreeNode* copyTree(TreeNode* other);
```

```

public:
    BST(); // Constructor
    ~BST(); // Destructor
    BST(const BST& other); // Copy Constructor
    BST& operator=(const BST& other); // Copy Assignment
    BST(BST&& other) noexcept; // Move Constructor
    BST& operator=(BST&& other) noexcept; // Move Assignment

    void insert(const T& value);
    std::vector<T> preOrderTraversal();
};

```

2. Implement TreeSort using the BST

✅ Steps to Implement TreeSort:

1. **Insert all elements** from an unsorted vector into the BST.
2. **Extract elements using pre-order traversal** and store them in another vector.
3. **Return the sorted vector.**

TreeSort Function:

```

template<typename T>
std::vector<T> treeSort(std::vector<T>& arr) {
    BST<T> tree;
    for (const T& val : arr) {
        tree.insert(val);
    }
    return tree.preOrderTraversal();
}

```

3. Compare Performance with std::sort

✅ Use the C++ <algorithm> library to compare execution time

- **Generate large random datasets (10,000+ elements)**
- **Measure sorting time for both methods using <chrono>**

Example Benchmarking Code:

```

int main() {
    const int SIZE = 10000;
    std::vector<int> data(SIZE);

    // Fill with random values
    for (int& val : data) {
        val = rand() % 10000;
    }

    // Measure TreeSort time
    auto treeData = data;
    auto start = std::chrono::high_resolution_clock::now();
    std::vector<int> treeSorted = treeSort(treeData);
    auto end = std::chrono::high_resolution_clock::now();
    std::cout << "TreeSort Time: "
              << std::chrono::duration<double, std::milli>(end - start).count()
              << " ms\n";

    // Measure std::sort time
    auto stdData = data;
    start = std::chrono::high_resolution_clock::now();
    std::sort(stdData.begin(), stdData.end());
    end = std::chrono::high_resolution_clock::now();
    std::cout << "std::sort Time: "
              << std::chrono::duration<double, std::milli>(end - start).count()

```

```

        << " ms\n";

    return 0;
}

```

Submission Requirements

- Submit a **C++ project (.cpp and .h files)** implementing:
 - **BST Class (with Big Five)**
 - **TreeSort Algorithm**
 - **Performance Comparison**
- Include a **README.md** with:
 - **Explanation of TreeSort and how it was implemented**
 - **Benchmark results comparing TreeSort vs. std::sort**
- Code should be **well-documented** and use **good object-oriented principles**.

Grading Criteria

Category	Points
BST correctly implemented with Big Five	30
TreeSort correctly implemented using BST	30
Pre-order traversal extracts sorted elements	20
Performance comparison with <code>std::sort</code>	10
Code readability, structure, and comments	10
Total	100

Bonus Challenge (+10 Points)

- Implement a **Self-Balancing BST (like AVL or Red-Black Tree)** and compare performance with the unbalanced BST for different inputs (sorted input, reverse-sorted input, randomized input).
- Analyze and **graph the performance difference** using different dataset sizes.

Bonus Challenge 2 (+10 Points)

- The main bottleneck for treesort is the repeated use of 'new' to allocate heap memory whenever a new node is created. To avoid the repeated use of 'new' (which results in an expensive syscall), modify your tree to do the following:
 - If the list to be sorted is very small (<1000 elements), **use stack memory** to store the nodes.
 - If the list is larger, use a custom allocator that calls 'new' only once to allocate a **memory pool** of the exact size needed to store all the nodes in the tree, and access the memory pool (instead of calling 'new' repeatedly) as nodes are added to the tree.

Expected Output Example

```

TreeSort Time: 35.2 ms
std::sort Time: 2.1 ms

```

👉 **std::sort should be significantly faster than TreeSort due to its optimized hybrid sorting approach.**

Appendix (hints and additional information):

In-order traversal (Left → Root → Right) prints the BST in sorted order.

Why?

- In a **Binary Search Tree (BST)**, the **left subtree** contains values smaller than the root.
- The **right subtree** contains values larger than the root.
- **Visiting nodes in the order Left → Root → Right** means:
 - **First, visit the smallest elements (leftmost) first.**
 - **Then, visit the root (middle value).**
 - **Finally, visit the right subtree (larger values).**
- This results in a **sorted sequence**.

In-order Traversal (Sorted Order)

✅ **Output:**

1 3 4 6 7 8 10 13 14


Pre-order Traversal (Root First)

❌ **Output (Not Sorted):**

8 3 1 6 4 7 10 14 13

♦ **Conclusion:**

If you want to **print the BST in sorted order**, use in-order traversal!



Hints for Implementing the **Big Five** in a Binary Search Tree (BST)

When implementing a **Binary Search Tree (BST)** in C++, you need to carefully manage **dynamic memory** to avoid memory leaks and undefined behavior. This is where the **Big Five** come into play:

1. Destructor (~BST()) – Freeing Allocated Memory

Hint:

- Since your BST **allocates nodes dynamically (new)**, you must **delete all nodes** in the destructor.
- **Use a recursive helper function to delete the entire tree** from bottom to top.

Guiding Question:

- How do you safely delete an entire tree **starting from the leaves**?

♦ **Think about a recursive post-order deletion method:**

```
void destroyTree(TreeNode* node) {
    if (node) {
        destroyTree(node->left);
        destroyTree(node->right);
        delete node;
    }
}
```

- Call this in the destructor.

2. Copy Constructor (BST(const BST& other))

Hint:

- If you create a new BST as a **copy of another BST**, you need to **deep copy** all nodes (not just copy pointers!).
- Use a **recursive function** to duplicate all nodes **from the original tree into the new tree**.

Guiding Question:

- How do you ensure that each node in the copied tree is a **newly allocated copy** of the corresponding node in the original tree?

♦ Think about a recursive copy function:

```
TreeNode* copyTree(TreeNode* otherNode) {
    if (!otherNode) return nullptr;
    TreeNode* newNode = new TreeNode(otherNode->data);
    newNode->left = copyTree(otherNode->left);
    newNode->right = copyTree(otherNode->right);
    return newNode;
}
```

- Call this in the copy constructor.
-

3. Copy Assignment Operator (BST& operator=(const BST& other))

Hint:

- The assignment operator needs to **copy data from an existing tree to an already existing tree**.
- Avoid **memory leaks** by first **deleting the old tree** before copying the new one.
- Handle **self-assignment** (if (this == &other) return *this;).

Guiding Question:

- What should you do first: **delete the current tree** or **copy the new tree**?

♦ Think about three steps:

1. **Check for self-assignment** (this != &other).
 2. **Destroy the existing tree** (to prevent memory leaks).
 3. **Copy the new tree recursively**.
-

4. Move Constructor (BST(BST&& other) noexcept)

Hint:

- When moving a BST, you want to **steal the pointer** to the other tree and **set the old tree to nullptr**.
- This avoids unnecessary deep copies and makes moving efficient.

Guiding Question:

- How do you **take ownership** of another tree without copying it?

♦ Think about what happens to the root pointer:

```
this->root = other.root;
other.root = nullptr; // Leave the moved-from tree empty
```

- No need to delete anything in this case.
-

5. Move Assignment Operator (BST& operator=(BST&& other) noexcept)

Hint:

- The move assignment is similar to the move constructor but also needs to **clear any existing tree data first**.
- Just like in the copy assignment operator, **handle self-assignment**.

Guiding Question:

- What do you do before stealing the other tree's pointer?

♦ Think about these steps:

1. **Check for self-assignment** (if (this == &other) return *this;).
 2. **Delete the existing tree**.
 3. **Steal the pointer** (this->root = other.root;).
 4. **Set other.root to nullptr** to prevent double deletion.
-

Final Thoughts

- **Write the Destructor First** – It will help you with the other functions.
- **Write a Recursive copyTree() Helper** – It simplifies both the copy constructor and copy assignment.
- **Think About Ownership** – When moving, transfer the pointer instead of copying nodes.

Once you implement these correctly, your BST will **avoid memory leaks** and **follow proper C++ ownership semantics**. 🚀